

CS51501 Spring 2016

Homework set # 2

Part (a): due at 11:59 pm on Monday Feb. 8, 2016
Part (b): due at 11:59 pm on Wednesday Feb. 17, 2016

Part (a)

Implement Algorithm 3.3 (BBTS: Block banded triangular solver)¹ on Page 60 in the textbook² to solve $\mathbf{L}\mathbf{x} = \mathbf{b}$, where \mathbf{L} is a banded lower triangular matrix with m subdiagonals.

1. Your code should be implemented using OpenMP and be tested on one node of the MC cluster.

2. In step 3 of BBTS, solve $\mathbf{L}_i(\mathbf{G}_i^{(0)}, \mathbf{f}_i^{(0)}) = (\mathbf{R}_i, \mathbf{f}_i)$ using the Column-sweep method in Algorithm 3.1(CSweep).

Download L.mtx and b.mtx on the course web page. Use your code to solve $\mathbf{L}\mathbf{x} = \mathbf{b}$, and store the solution \mathbf{x} to x.mtx. The order of \mathbf{L} is $n = 8192$, with $m = 128$.

TA will test your code as follows: `./bbts A.mtx b.mtx x.mtx`.

Part (b)

Implement a SPIKE-like algorithm to solve $\hat{\mathbf{L}}\mathbf{y} = \mathbf{f}$, where $\hat{\mathbf{L}}$ is a banded lower triangular matrix of order $N = 32,768 = 2^{15}$, with $t = 128$ subdiagonals on a cluster of multicore nodes. Your code should be implemented using the MPI/OpenMP programming paradigm, and be tested on four nodes of the MC cluster. The SPIKE-like algorithm should be implemented using MPI. In this algorithm, use BBTS (from Part (a)) to solve the system involvings $\hat{\mathbf{L}}_i$.

1. Partition $\hat{\mathbf{L}}$ and \mathbf{f} into 4 parts, with one partition in each node, via `MPI_Scatter()`. Therefore, node i should contain $\hat{\mathbf{L}}_i$ and \mathbf{R}_i , $i = 1 : 4$. Note that

¹Correction(BBTS):

step 2: `doall 2: n/m`

step 3: solve $\mathbf{L}_i(\mathbf{G}_{i-1}^{(0)}, \mathbf{f}_i^{(0)}) = (\mathbf{R}_{i-1}, \mathbf{f}_i)$

²Gallopoulos, Efstratios, Bernard Philippe, and Ahmed H. Sameh. Parallelism in Matrix Computations. Springer, 2015.

$\hat{\mathbf{L}}_i$ is the banded lower triangular matrix shown below, with \mathbf{R}_i being an upper triangular matrix.

$$\left[\begin{array}{c|c|c|c} \hat{\mathbf{L}}_1 & & & \\ \hline & \mathbf{R}_2 & \hat{\mathbf{L}}_2 & \\ & \mathbf{0} & & \\ \hline & & \mathbf{R}_3 & \hat{\mathbf{L}}_3 \\ & & \mathbf{0} & \\ \hline & & & \mathbf{R}_4 & \hat{\mathbf{L}}_4 \\ & & & \mathbf{0} & \end{array} \right] \left[\begin{array}{c} \mathbf{y}_1 \\ \hline \mathbf{y}_2 \\ \hline \mathbf{y}_3 \\ \hline \mathbf{y}_4 \end{array} \right] = \left[\begin{array}{c} \mathbf{f}_1 \\ \hline \mathbf{f}_2 \\ \hline \mathbf{f}_3 \\ \hline \mathbf{f}_4 \end{array} \right].$$

2. Solve

$$\begin{aligned} \mathbf{L}_1 \begin{pmatrix} \mathbf{g}_1 \\ \mathbf{h}_1 \end{pmatrix} &= \mathbf{f}_1, \\ \mathbf{L}_i \left[\begin{pmatrix} \mathbf{g}_i \\ \mathbf{h}_i \end{pmatrix}, \begin{pmatrix} \mathbf{M}_i \\ \mathbf{N}_i \end{pmatrix} \right] &= \left[\mathbf{f}_i, \begin{pmatrix} \mathbf{R}_i \\ \mathbf{0} \end{pmatrix} \right], \end{aligned}$$

where $i=2:4$, via BBTS (from Part (a)). Note that the square matrices \mathbf{N}_i are of order t , and the vector \mathbf{h}_i are of order t .

Therefor, as you have seen in the lecture, the resulting system is of the form

$$\left[\begin{array}{c|c|c|c} \mathbf{I} & & & \\ \hline & \mathbf{M}_2 & \mathbf{I} & \\ & \mathbf{N}_2 & & \\ \hline & & \mathbf{M}_3 & \mathbf{I} \\ & & \mathbf{N}_3 & \\ \hline & & & \mathbf{M}_4 & \mathbf{I} \\ & & & \mathbf{N}_4 & \end{array} \right] \left[\begin{array}{c} \mathbf{y}_1 \\ \hline \mathbf{y}_2 \\ \hline \mathbf{y}_3 \\ \hline \mathbf{y}_4 \end{array} \right] = \left[\begin{array}{c} \mathbf{g}_1 \\ \mathbf{h}_1 \\ \hline \mathbf{g}_2 \\ \mathbf{h}_2 \\ \hline \mathbf{g}_3 \\ \mathbf{h}_3 \\ \hline \mathbf{g}_4 \\ \mathbf{h}_4 \end{array} \right].$$

3. In node 1, gather \mathbf{N}_i and \mathbf{h}_i from node i , $i = 2, 3$ and 4 via MPI_Gather(). Solve the reduced system

$$\left[\begin{array}{cccc} \mathbf{I} & & & \\ \mathbf{N}_2 & \mathbf{I} & & \\ & \mathbf{N}_3 & \mathbf{I} & \\ & & \mathbf{N}_4 & \mathbf{I} \end{array} \right] \left[\begin{array}{c} \mathbf{u}_1 \\ \mathbf{u}_2 \\ \mathbf{u}_3 \\ \mathbf{u}_4 \end{array} \right] = \left[\begin{array}{c} \mathbf{h}_1 \\ \mathbf{h}_2 \\ \mathbf{h}_3 \\ \mathbf{h}_4 \end{array} \right]$$

via CSweep.

4. Send \mathbf{u}_1 to node 2, \mathbf{u}_2 to node 3, \mathbf{u}_3 to node 4, and simultaneously compute

$$\begin{aligned} \mathbf{y}_1 &= \begin{pmatrix} \mathbf{g}_1 \\ \mathbf{h}_1 \end{pmatrix}, \\ \mathbf{y}_i &= \begin{pmatrix} \mathbf{g}_i \\ \mathbf{h}_i \end{pmatrix} - \begin{pmatrix} \mathbf{M}_i \\ \mathbf{N}_i \end{pmatrix} \mathbf{u}_{i-1}, \end{aligned}$$

where $i=2:4$

5. Gather \mathbf{y}_i via `MPI_Gather()` and write it to `y.mtx`.

Download `L.hat.mtx` and `f.mtx` on the course web page. Use your code to solve $\hat{\mathbf{L}}\mathbf{y} = \mathbf{f}$, and store the solution \mathbf{y} to `y.mtx`.

TA will test your code as follows:

```
mpirun -np 4 -hostfile nodes ./spike.like L.hat.mtx f.mtx y.mtx
```

Storage

1. vectors will be stored as dense in `mtx` format. Reading \mathbf{b} and \mathbf{f} is the same as HW1.

2. Storing a banded lower triangular matrix in dense `mtx` format will consume a lot of memory needlessly. So TA stores it as a sparse matrix in coordinate format. After you read the matrix from the `mtx` file, you should store it properly. See band storage on <https://software.intel.com/en-us/node/520871>. Please do not store the matrix as a 2-dimensional array of order n .

The Intel compiler

From HW2, you need to use the Intel compiler(i.e. `icc`, `mpiicc`, `ifort`, `mpiifort`) to compile your code. In later HW sets, you will learn how to use the Intel MKL library.

0.1 Set up environment for the Intel compiler

TA prepared a `defs` file to set up the environment for the Intel compiler.

1. Download `mc.defs` on the course web page and put this file in your home directory.

2. After you access `mcxx.cs.purdue.edu`, type `'source mc.defs'`.

Then you could use `icc`, `mpiicc`, `ifort` and `mpiifort`.

0.2 Commands to compile the OpenMP code

C/C++: `icc -openmp test.c -o test`

FORTTRAN: `ifort -openmp test.f -o test`

Note that the compile-time flag `-fopenmp` is changed to `-openmp`.

Submission

Part (a): `turnin -c cs51501 -p HW2a your_folder_name`

Part (b): `turnin -c cs51501 -p HW2b your_folder_name`

Your submission should include the following files:

1. The source code.
2. A Readme file or a Makefile, which includes all the compiling commands.

Note that the submission of Part (b) should include the BBTS function from Part (a).

Please **do not** include the test cases in your submission, which would be too large to submit.